

1 Introduction

Taking notes from the first recorded lecture. The speaker says that this is about what to expect from the *second part of this course*.

Most modern processor architectures are multicore - threads can be executed in parallel. Other than architecture level concurrency, distributed computing systems, e.g. databases, involve concurrent execution so that any one node can provide a coherent and uniform picture of the database irrespective of where you are in the world.

Support for concurrency

- architecture level
- distributed systems
- programming languages (modern standards of C++, Java etc, modern languages like Rust)

Naming conventions

- Shared variables: **x**, **y**
- Local variables/registers: **a**, **b**, **c**, **d**

Consider two threads

Table 1: initially **x** = 0

th_1	th_2
x :=1	x :=2
a := x	b := x

Traditional semantics used to implement concurrent programs - sequential concurrency (SC). From what I understood, this means that two threads cannot be interleaved; execution of one thread cannot be interrupted by a "context switch". You always read from the latest possible write. It was summarized in the lecture as

- Shared memory
- Processes/threads: atomic read/write
- interleaving of operations across threads
- simple, but too strong

Too strong? If you look at modern system designs - microprocessors, cache protocols, distributed computing don't use SC, but use weak memory models instead.

- Intel: TSO
- IBM: Power
- ARM

As for programming languages

- C11
- Java
- Rust

provide first-class support for concurrency use release-acquire semantics, which is also a weak memory model. The behaviour of your program under this semantics will be different from what you can expect for sequential consistency.

Consider the example of a simple program. Init: $x=y=0$.

Table 2: two threads

$x=1;$	$y=1;$
$ry=y;$	$rx=x;$

[Second sitting begins here. Will not try to make notes for everything that the speaker says - only handful of important points]

In the above threads, using SC cannot yield a state where both rx and ry are 0. On the other hand, if we allow "context switches" between a thread execution, i.e. under release-acquire semantics, we can reach a state where $rx=ry=0$.

Different memory models allow different constraints for how threads can be interleaved. General class - weak memory models.

In RA semantics - release writes, acquire reads (look up on your fave search engine).

Reachability problem: Consider a general multithreaded program with a bunch of shared resources and local variables. The question we are interested in - *is this program unsafe?* Is it possible to reach a bad state? (Under some semantics). Can we write an algo that can tell whether such a bad state is reachable?

So, we want to build a tool that takes a concurrent program as an input. The problem could be either decidable or undecidable. If it the reachability is determinable, then there'll be some complexity and all that shebang.

The other aspect is ensuring that a bad state can never be reached. Of course, we cannot really go about exploring all possible reachable states. There should be a way of "pruning the search".

So, a process π is a sequence of instructions. There is another abstraction - trace τ . A trace is essentially a graph. Each node is an atomic instruction. Edges can be

- **po** - program order
- **rf** - read-from
- **co** - coherence order
- **fr** - from-read

SC is guaranteed if the graph is acyclic. We say that $\pi \models \tau$ (read as "pi models tau").

Once you have a particular trace, you can look at different processes π_i to see whether they are consistent with the trace.

We explore the program by generating and exploring traces. Sufficient to explore one run per trace to verify. In general multiple runs correspond to a single trace.




Figure 1: Example of exploring using traces

2 Some Self Study

This section is comprised of notes taken while reading from Shavit et al.

Gotta start thinking like a concurrent programmer - understanding when operations "happen", identifying all possible interleavings etc.

Two complementary directions

- Principles - computability, what can be computed in an asynchronous concurrent environment
- Practice

Consider the example of finding all prime numbers $< 10^{10}$, and you have ten threads at your disposal. A naive approach would be to distribute the input domain between all threads equally, but this would be far from optimum given how load balancing is affected.

An alternative would be to use a shared counter and assign an integer to each thread.

```
// a naive implementation of shared counter
public class Counter {
    private long value;
    public Counter(long i) {
        value = i;
    }
    public long getAndIncrement() {
```

```

        return value++;
    }
}

Counter counter = new Counter(1);
void primePrint {
    long i = 0;
    long limit = power(10,10);
    while (i < limit) {
        i = counter.getAndIncrement();
        if (isPrime(i))
            print(i);
    }
}

```

Note that the above implementation of the shared counter is a bit naive, because the instruction `return value++;` is actually something like

```

long temp = value;
value = temp + 1;
return temp;
value is shared. temp is local.

```

Modern microprocessors have hardware support for atomic instructions that update shared counters etc (mutex?). Alternatively a software solution would require ensuring that only one thread updates the shared variable before another thread can do anything with it. E.g. Rust compilers should be able to do this statically. This is a classic problem in multiprocessing algorithms, called the *mutual exclusion problem*.

I've heard some of these terms before

- mutex
- deadlock
- bounded fairness
- blocking vs non-blocking synchronization

3 First Lecture

The course is about software verification, specifically concurrent software. Very open problem. Situations can be decidable or undecidable.

The technology is getting there, but it's not *there*-there. If you want to verify that your software, you know, works, it's not as straightforward as just downloading a verification tool and reading its API documentation. You need a whole team that understands the whole verification shebang.

Even if you have only two threads, the situation can become significantly more complicated than a similar sequential situation. Rust was made because most bugs on Firefox were concurrency bugs.

Question: What is the number of schedules between two threads with number of instructions N_1 and N_2 ?

$$^{N_1+N_2}C_{N_1}.$$

