

CS 213 - Data Structures and Algorithms

January 6, 2021

1 Introduction

Twitter's banner problem - the agenda is to analyze this problem and work towards a solution. (We were supposed to see a video - I missed it lmao)

CEO of twitter - Parag Agarwal - his student. Well anyway. What's the problem at hand? People post content on twitter. There's some four people on the banner page of twitter. Those four people will show up to scrapers etc.

Problem: Place "most popular posts on the top page dynamically - updated every five-ish minutes or so."

If some folk says something super-interesting that has the potential to go super-viral, then you want to shove that on the banner page.

So, consider an array of popularity indices. You as a programmer need to process this. We will work with toy models only - no full sized user bases and all.

We want to get the most popular from the list. $\mathcal{O}(n)$ to get **max**, then there's going to be a hole to plug. One approach - sort the list, pop from the stack. This way no holes to be dealt with. Problem with sorting (in this particular situation) - if a super-interesting post comes in - we will not be in a position to dynamically adjust the stack to prioritize it. Also, sorting is something like $\mathcal{O}(n^2)$ - number of cycles doubles in order of magnitude.




Figure 1: figures/btree.png

The data can be reorganized into another data structure that is more convenient w.r.t insertions, searching etc. See the above figure. That is a complete,

ordered binary tree (ordered \implies left child and right child are distinct for each node, binary \implies two leaves per node, 2^i nodes in level i).

Internally, a binary tree such as this is still stored as an array. If we want the right child of j , we only need realize $j \rightarrow 2j + 1$., which is like ultra easy for the computer.

- Right child: $j \rightarrow 2j + 1$
- Parent: $k \rightarrow \frac{k}{2}$

The tree has a

- Structure property: it's a complete B tree
- Comparison property: the internal array is sorted - reflects in the B Tree representation

So, how would we go about handling the banner problem? The problem is finding the most popular, next most popular and so on... And the twist is that there are insertions and deletions happening dynamically in the internal array.

So, the first part is *Extract max*. Consider the figure above. This is what

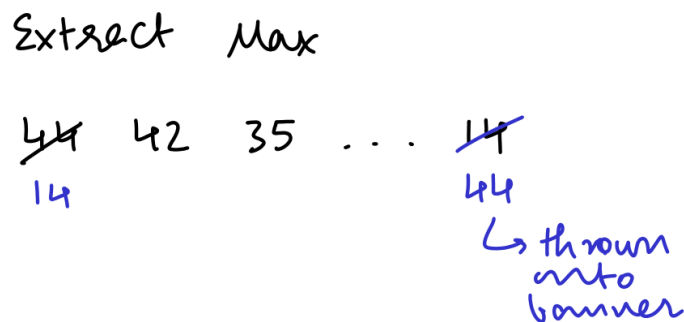


Figure 2: figures/extractmax.png

we're doing when we select the maximum from the array. 44 and 14 have been mutually swapped. But, this destroys the ordering property.

We can then perform a sequence of operations to fix the order. We can keep




Figure 3: figures/fix.png

performing swaps between parent and child as long as we maintain the invariant that **parent** > **child**. And because. There will be at most $\log_2 n$ comparisons required, where n is the number of elements in the internal array.

Now, if there is some database trigger and the stack gets updated, then in general the comparison property will be destroyed. We can do the same process of sorting them - this time propagating up or down depending on how the comparisons hold up.




Figure 4: figures/insertion.png

Heap: a complete binary tree with the comparison property. Selecting the right data structure is like choosing a screddriver for a screw that you want to work with. You'd want to choose the best fit, ya? Don't want to be shoving square cocks in triangle holes.